

In [1]:

```
%matplotlib inline

import sys
import os
import sklearn
import numpy as np
import scipy as sp
import scipy.sparse as spa
import matplotlib.pyplot as plt
from time import time
from sklearn.datasets import fetch_20newsgroups
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.naive_bayes import MultinomialNB
from sklearn.cluster import KMeans
```

In [2]:

```
# Load the training set
print("Loading 20 newsgroups training set... ")
dataset = fetch_20newsgroups(subset='all', shuffle=True, random_state=4)
#Classifying dataset into training set
t0=time()
newsgroups_train = fetch_20newsgroups(subset='train')
#print(newsgroups_train.DESCR)
print("%d documents" % len(newsgroups_train.filenames))
#filenames = newsgroups_train.filenames
print("%d categories" % len(newsgroups_train.target_names))
print("done in %fs" % (time() - t0))
```

```
Loading 20 newsgroups training set...
11314 documents
20 categories
done in 0.285633s
```

In [3]:

```
print("Loading 20 newsgroups test set... ")
#Classifying dataset into test set
newsgroups_test = fetch_20newsgroups(subset='test')
t0 = time()
print("done in %fs" % (time() - t0))
print("Predicting the labels of the test set...")
print("%d documents" % len(newsgroups_test.filenames))
print("%d categories" % len(newsgroups_test.target_names))
```

```
Loading 20 newsgroups test set...
done in 0.000000s
Predicting the labels of the test set...
7532 documents
20 categories
```

In [4]:

```
# Feature Extraction  
#We now vectorize the dataset using SKLearn vectorizers.  
#This is the stage of feature extraction. We use the TF-IDF feature model.  
t0=time()  
vectorizer = TfidfVectorizer(encoding='latin1')  
X_train = vectorizer.fit_transform(newsgroups_train.filenames)  
print("done in %fs" % (time() - t0))  
print("n_samples: %d, n_features: %d" % X_train.shape)  
assert spa.issparse(X_train)  
y_train = newsgroups_train.target
```

```
done in 0.221436s  
n_samples: 11314, n_features: 9879
```

In [6]:

```
print("Extracting features from the dataset using the same vectorizer")  
t0 = time()  
X_test = vectorizer.transform( newsgroups_test.filenames)  
y_test = newsgroups_test.target  
print("done in %fs" % (time() - t0))  
print("n_samples: %d, n_features: %d" % X_test.shape)
```

```
Extracting features from the dataset using the same vectorizer  
done in 0.157541s  
n_samples: 7532, n_features: 9879
```

In [5]:

```
#Benchmarking aims at evaluating something by comparison with a standard.  
#Benchmarking the code means how fast the code is executing and where the bottleneck is.  
#One major reason for benchmarking is that it optimizes the code.  
  
#A confusion matrix is a tabular summary of the number of correct and incorrect predictions made by a classifier.  
#It can be used to evaluate the performance of a classification model through the calculation of performance metrics like accuracy, precision, recall, and F1-score.
```

In [6]:

```
#benchmark MultinomialClassifiers

def benchmark(clf_class, params, name):
    print("parameters:", params)
    t0 = time()
    clf = clf_class(**params).fit(X_train, y_train)
    print("done in %fs" % (time() - t0))

    if hasattr(clf, 'coef_'):
        print("Percentage of non zeros coef: %f" % (np.mean(clf.coef_ != 0) * 100))
    print("Predicting the outcomes of the testing set")
    t0 = time()
    pred = clf.predict(X_test)
    print("done in %fs" % (time() - t0))

    print("Classification report on test set for classifier:")
    print(clf)
    print()
    print(classification_report(y_test, pred, target_names=newsgroups_test.target_names
))

    cm = confusion_matrix(y_test, pred)
    print("Confusion matrix:")
    print(cm)

# Show confusion matrix
plt.matshow(cm)
plt.colorbar(orientation = 'horizontal' )
plt.title('Confusion matrix')
```

In [9]:

```
#Benchmark Multinomial Classifier  
print("Testbenching a MultinomialNB classifier...")  
parameters = {'alpha': 0.01}  
  
benchmark(MultinomialNB, parameters, 'MultinomialNB')
```

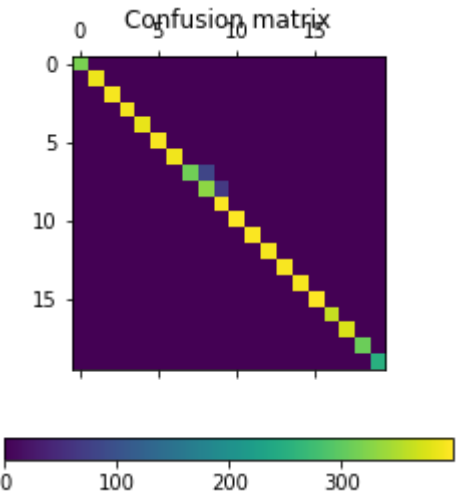
Testbenching a MultinomialNB classifier...
parameters: {'alpha': 0.01}
done in 0.021942s
Percentage of non zeros coef: 100.000000
Predicting the outcomes of the testing set
done in 0.003990s
Classification report on test set for classifier:
MultinomialNB(alpha=0.01, class_prior=None, fit_prior=True)

	precision	recall	f1-score	support
alt.atheism	1.00	1.00	1.00	319
comp.graphics	1.00	1.00	1.00	389
comp.os.ms-windows.misc	1.00	1.00	1.00	394
comp.sys.ibm.pc.hardware	1.00	1.00	1.00	392
comp.sys.mac.hardware	1.00	1.00	1.00	385
comp.windows.x	1.00	1.00	1.00	395
misc.forsale	1.00	1.00	1.00	390
rec.autos	1.00	0.79	0.88	396
rec.motorcycles	0.80	0.83	0.81	398
rec.sport.baseball	0.86	1.00	0.92	397
rec.sport.hockey	1.00	1.00	1.00	399
sci.crypt	1.00	1.00	1.00	396
sci.electronics	1.00	1.00	1.00	393
sci.med	1.00	1.00	1.00	396
sci.space	1.00	1.00	1.00	394
soc.religion.christian	1.00	1.00	1.00	398
talk.politics.guns	1.00	1.00	1.00	364
talk.politics.mideast	1.00	1.00	1.00	376
talk.politics.misc	1.00	1.00	1.00	310
talk.religion.misc	1.00	1.00	1.00	251
accuracy			0.98	7532
macro avg	0.98	0.98	0.98	7532
weighted avg	0.98	0.98	0.98	7532

Confusion matrix:

[[319	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0]																
[0	389	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0]																
[0	0	394	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0]																
[0	0	0	392	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0]																
[0	0	0	0	385	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0]																
[0	0	0	0	0	395	0	0	0	0	0	0	0	0	0	0	0	0
0	0]																
[0	0	0	0	0	0	0	312	84	0	0	0	0	0	0	0	0	0
0	0]																
[0	0	0	0	0	0	0	0	331	67	0	0	0	0	0	0	0	0
0	0]																
[0	0	0	0	0	0	0	0	0	397	0	0	0	0	0	0	0	0
0	0]																
[0	0	0	0	0	0	0	0	0	0	399	0	0	0	0	0	0	0
0	0]																
[0	0	0	0	0	0	0	0	0	0	0	396	0	0	0	0	0	0
0	0]																

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0 393  0  0  0  0  0
 0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0 395  1  0  0  0  0
 0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0 394  0  0  0  0
 0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0 398  0  0  0
 0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0 364  0  0
 0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0 376  0
 0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
310 0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0 251]]
```



In [7]:

```
#removing the noisy data
```

```
vectorizer = TfidfVectorizer(min_df=10, max_df=0.5, stop_words='english')  
vectorized = vectorizer.fit_transform(dataset.data)  
num_samples, num_features = vectorized.shape  
num_clusters = 20
```

```
#creating clusters
```

```
km = KMeans(n_clusters=num_clusters, init='random', max_iter=30, n_init=1, verbose=0 )  
km.fit(vectorized)  
labels = np.array(km.labels_)
```

```
#printing clusters
```

```
for i in range(0, num_clusters):  
    num_docs_in_cluster = sum((labels == i))  
    print ("cluster ", i, ": ", "number of documents ", num_docs_in_cluster)
```

```
cluster 0 : number of documents 3801  
cluster 1 : number of documents 128  
cluster 2 : number of documents 648  
cluster 3 : number of documents 133  
cluster 4 : number of documents 538  
cluster 5 : number of documents 123  
cluster 6 : number of documents 4958  
cluster 7 : number of documents 2800  
cluster 8 : number of documents 621  
cluster 9 : number of documents 402  
cluster 10 : number of documents 249  
cluster 11 : number of documents 550  
cluster 12 : number of documents 555  
cluster 13 : number of documents 241  
cluster 14 : number of documents 885  
cluster 15 : number of documents 390  
cluster 16 : number of documents 1176  
cluster 17 : number of documents 184  
cluster 18 : number of documents 270  
cluster 19 : number of documents 194
```

In [8]:

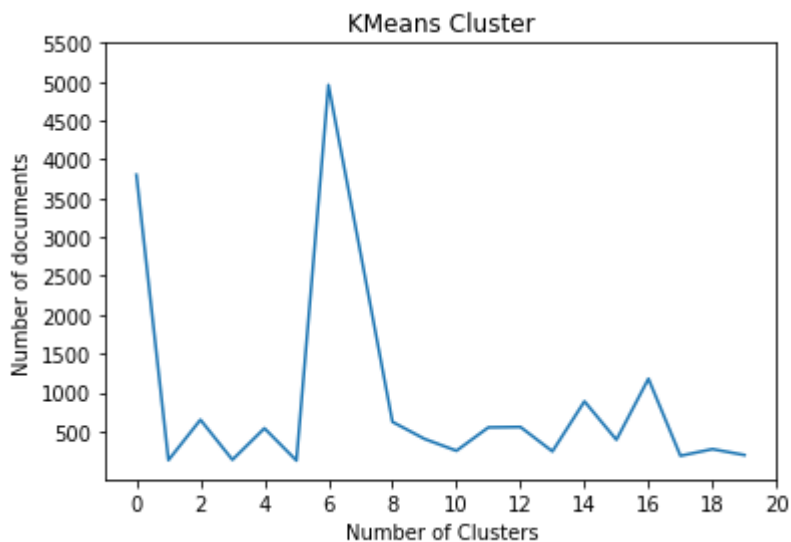
```

arr_1 = [0 for x in range(20)]
for x in range (0, 20):
    arr_1[x]=x

arr_2=[0 for y in range(20)]
for y in range(0,num_clusters):
    num_docs_in_cluster = sum((labels == y))
    arr_2[y]=num_docs_in_cluster

#Line Plot
plt.plot(arr_1,arr_2)
plt.xlabel('Number of Clusters')
plt.ylabel ('Number of documents')
plt.title('KMeans Cluster')
plt.xticks((0,2,4,6,8,10,12,14,16,18,20))
plt.yticks((500,1000,1500,2000,2500,3000,3500,4000,4500,5000,5500))
plt.show()

```



In [9]:

```

# Input post
new_post=input("Enter new_post: ")
t0 = time()

```

Enter new_post: Disk drive problems. Hi, I have a problem with my hard disk. After 1 year it is working only sporadically now. I tried to format it, but now it doesn't boot any more. Any ideas? Thanks.

In [29]:

```
#finding similar posts
if new_post.isdigit():
    print("Invalid post")
else:
    # vectorizing input post
    new_post_vec = vectorizer.transform([new_post])
    new_post_label = km.predict(new_post_vec)[0]
    #Fetching indices in the original dataset
    similar_indices = (labels==new_post_label).nonzero()[0]
    #Finding similar posts
    similar = []
    for i in similar_indices:
        dist = sp.linalg.norm((new_post_vec-vectorized[i,:]).toarray())
        similar.append((dist, dataset.data[i]))
    similar = sorted(similar)

#print(similar)

print("Number of similar posts: ", len(similar))
print("done in %fs" % (time() - t0))
```

Number of similar posts: 1176
done in 3391.058196s

In [28]:

```
# Number of documents in a cluster
docs = [0 for i in range(20)]
for i in range(0,num_clusters):
    num_docs_in_cluster = sum((labels == i))
    docs[i]= num_docs_in_cluster

temp = np.append(np.cumsum(docs[:-1]),docs[-1])
temp[19]= temp[18]+temp[19]

#finding similar posts in each cluster
x = 0
similar_posts_in_each_cluster = [0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0]
for i in range(0,20):
    for j in range(0, len(similar_indices)):
        if(similar_indices[j] < temp[i]):
            x = x+1

    similar_posts_in_each_cluster[i] = x


    x = 0
similar_posts_in_each_cluster
```

Out[28]:

```
[245,
 253,
 292,
 296,
 329,
 340,
 649,
 818,
 857,
 882,
 901,
 942,
 973,
 989,
 1044,
 1065,
 1139,
 1150,
 1167,
 1176]
```

In [27]:

```
#Line plot
plt.plot(arr_1, similar_posts_in_each_cluster,color='green', linestyle='dashed', linewidth = 3,
         marker='o', markerfacecolor='blue', markersize=12)
plt.xlim(0,20)
plt.ylim(0,1300)
plt.xlabel('Cluster Number')
plt.ylabel('Number of similar posts')
plt.title('Similar Posts in the Cluster')
plt.xticks((0,2,4,6,8,10,12,14,16,18,20))
plt.yticks((100,200,300,400,500,600,700,800,900,1000,1100,1200,1300))
plt.show()
```



In [20]:

```
show_at_1 = similar[0]
show_at_2 = similar[int(len(similar)/2)]
show_at_3 = similar[-1]

from tabulate import tabulate
table = [show_at_1, show_at_2, show_at_3]
print(tabulate(table, headers=["Index", "Position", "Excerpt Posts"], showindex=True, tablefmt="grid"))
```

```

+-----+-----+-----+
+-----+
|   Index |   Position | Excerpt Posts
|
+=====+=====+=====+
=====+
|         0 |    1.08548 | From: rogntorb@idt.unit.no (Torbj|rn Rognes)
|         |             | Subject: Adding int. hard disk drive to IICx
|         |             | Keywords: Mac IICx, internal, hard disk drive, SC
SI         |             |
|         |             | Reply-To: rogntorb@idt.unit.no (Torbj|rn Rognes)
|         |             | Organization: Div. of CS & Telematics, Norwegian
Institute of Technology |
|         |             | Lines: 32
|         |             |
|         |             |
|         |             | I haven't seen much info about how to add an extr
a internal disk to a
|         |             | mac. We would like to try it, and I wonder if som
eone had some good
|         |             | advice.
|         |             |
|         |             |
|         |             | We have a Mac IICx with the original internal Qua
ntum 40MB hard disk,
|         |             | and an unusable floppy drive. We also have a new
spare Connor 40MB
|         |             | disk which we would like to use. The idea is to r
eplace the broken
|         |             | floppy drive with the new hard disk, but there se
ems to be some
|         |             | problems:
|         |             |
|         |             |
|         |             | The internal SCSI cable and power cable inside th
e cx has only
|         |             | connectors for one single hard disk drive.
|         |             |
|         |             |
|         |             | If I made a ribbon cable and a power cable with t
hree connectors each
|         |             | (1 for motherboard, 1 for each of the 2 disks), w
ould it work?
|         |             |
|         |             |
|         |             | Is the IICx able to supply the extra power to the
extra disk?
|         |             |
|         |             |
|         |             | What about terminators? I suppose that i should r
emove the resistor
|         |             | packs from the disk that is closest to the mother
board, but leave them
|         |             | installed in the other disk.

```

```

|                                     |
|                                     |
| t the new disk gets               |
|                                     |
|                                     |
|                                     |
| e, as we have an                  |
|                                     |
| of the hard disk.                 |
|                                     |
|                                     |
|                                     |
|-----|
|                                     |
| rognatorb@idt.unit.no             |
| +-----+-----|
|-----|
|          1 |          1.39408      |
| igler)     |                       |
|                                     |
|                                     |
| eveland, OH (USA)                 |
|                                     |
| senbigler)                          |
|                                     |
|                                     |
|                                     |
|                                     |
| e and mother board problem        |
| iver for windows                   |
| problem windows would              |
| LOGO and either it would          |
| sistent EVERY OTHER TIME          |
|                                     |
| es the same thing.                 |
| t it worked FINE!                  |

```

```

|           |           | I can get into windows each time now with the w
in/s command.           |
|           |           | This forces Standard mode. Things seem to run sl
ower. I mainly           |
|           |           | use windows apps, but in standard mode there is n
o virtual mem.....    |
|           |           | plus it is slower.
|           |           |
|           |           |
|           |           | I re-loaded windows, it still does the same thin
g. Should I first delete |
|           |           | everything in all windows dir's? I did not becau
se I have so much        |
|           |           | added in sub dir's etc.
|           |           |
|           |           |
|           |           | Really puzzling why ENHANCED MODE would not load
each time but consistently |
|           |           | every other time. Standard mode each time.....
|           |           |
|           |           |
|           |           | ANY THOUGHTS OR COMMENTS ???
|           |           |
|           |           |
|           |           | C-ya..... /\ /\artin
|           |           | --
|           |           |
|           |           | This communication is sent by /\ /\artin Un
iversity of Arizona Tucson |
|           |           | =====
|           |           | =====
|           |           | ak333@cleveland.freenet.edu mlinseb@ccit.ari
zona.edu mlinseb@arizvms   |
|           |           | DEATH HAS BEEN DEAD FOR ABOUT 2,000 YEARS *****
* FOLLOW THE KING OF KINGS |
+-----+-----+-----+-----+-----+-----+
+-----+
|           | 2 | 1.41421 | From: dplatt@ntg.com (Dave Platt)
|           |   |       |
|           |   |       | Subject: Jumper settings for Ungermann-Bass PCNIC
Ethernet card           |
|           |   |       | Organization: New Technologies Group
|           |   |       |
|           |   |       | Distribution: usa
|           |   |       |
|           |   |       | Lines: 16
|           |   |       |
|           |   |       |
|           |   |       | Does anybody have a data-sheet handy for the abo
ve-mentioned card? I     |
|           |   |       | bought one, sans manual at a local surplus shop,
and want to try it out    |
|           |   |       | with the Crywyr packet driver suite.
|           |   |       |
|           |   |       |

```

```

|                                     | The IRQ and interface-select jumpers are pretty s
| straightforward, but I            |
| don't grok the settings of W10-W18 (also labelled
A15 through A18).
|                                     |
| Could somebody tell me which settings of these fo
ur jumpers correspond
|                                     |
| to what I/O addresses?
|                                     |
|                                     |
|                                     |
| Is there anything else about this card I should k
now, before I
|                                     |
| plug&play?
|                                     |
|                                     |
| --
|                                     |
| Dave Platt
VOICE: (415) 813-8917 |
|                                     | Domain: dplatt@ntg.com      UUCP:
...netcomsv!ntg!dplatt
|                                     |
| USNAI: New Technologies Group Inc. 2470 Embarca
rdero Way, Palo Alto CA 94303
+-----+
+-----+

```

In []: